

Regularization in RL

- DRA's KL
- [Todorov's KL](#)
 - Abstracts away from actions
 - KL on divergence term is of control dynamics from default dynamics
- [Linear RL](#)
 - RL version of Todorov's work where dynamics are replaced by policies
- [SAC](#)/maxEnt RL
- Moskovitz+Botvinick's regularization
 - unpublished
 - has two regularizing terms:
 1. the first term is the same as Linear RL, a KL between behavioral policy and default policy, which is learned here unlike linear RL/Todorov
 2. the second term is a complexity (max description length) penalty on the default policy
 - has an interesting information-theoretic story and easy implementation thanks to the legendary Kingma: both these terms are what come out of [variational dropout](#)
- [Noisy Nets](#)
 - Noisy weights in a DQN
 - maybe only related to DRA and none of the others

Ideas linked to DRA vs MaxEnt

- Read [lilian weng's blog on policy gradient algorithms](#)
 - DPG
 - DDPG
 - TRPO
 - PPO
 - SAC
- Read [noisynets](#) and [plappert et al 2017](#)